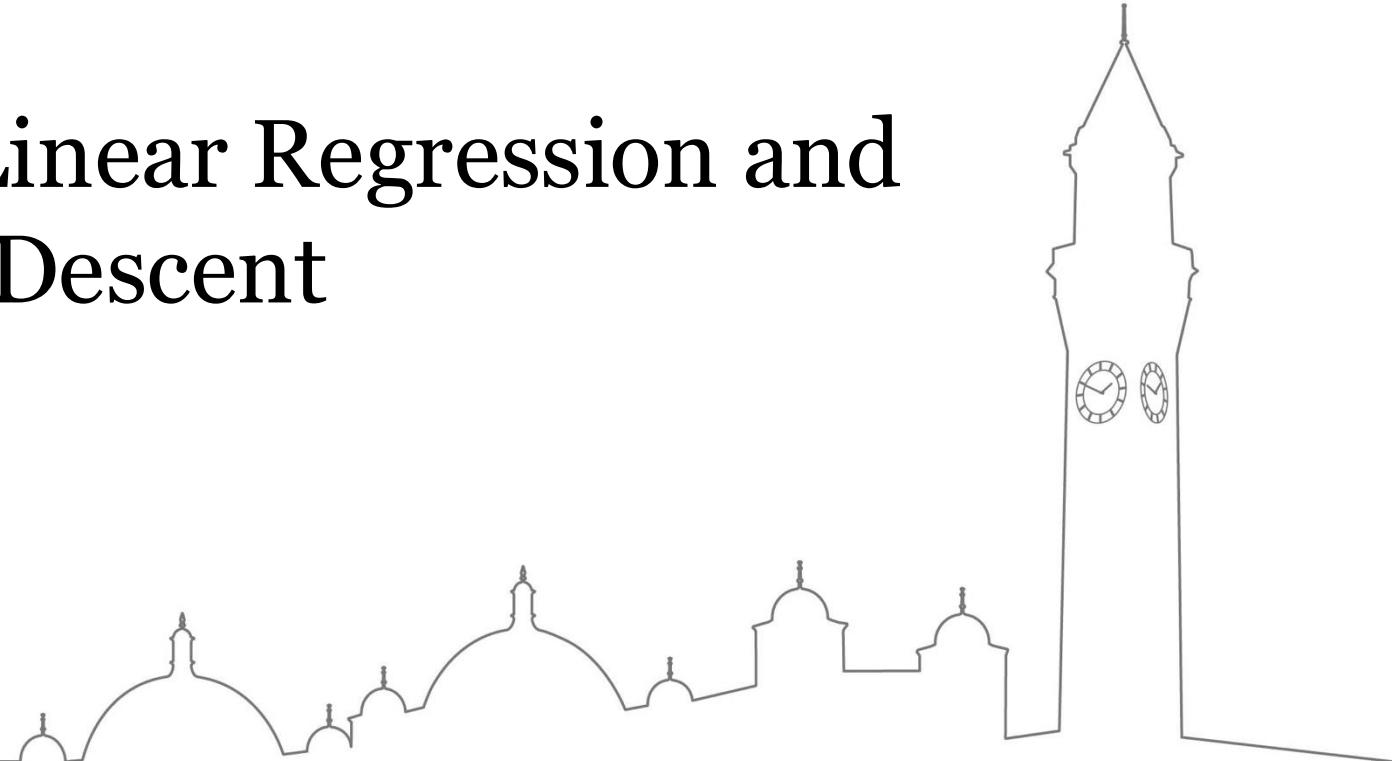




UNIVERSITY OF  
BIRMINGHAM

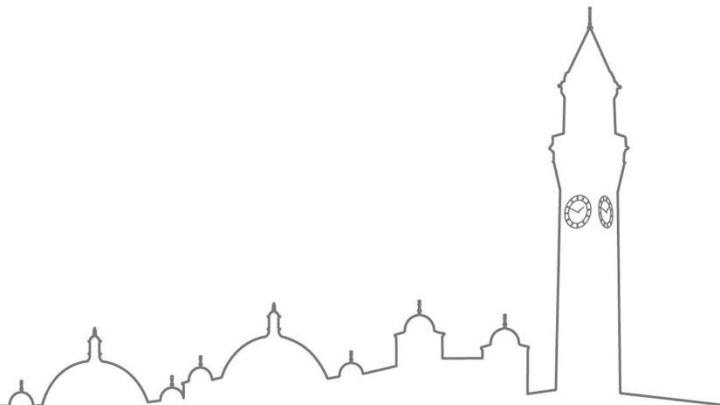
# Week 2. Linear Regression and Gradient Descent

Dr. Shuo Wang



# Overview

- Linear Regression – a ML algorithm for regression problems
- Gradient descent – an optimisation technique used to in ML algorithms.



# Recall: regression

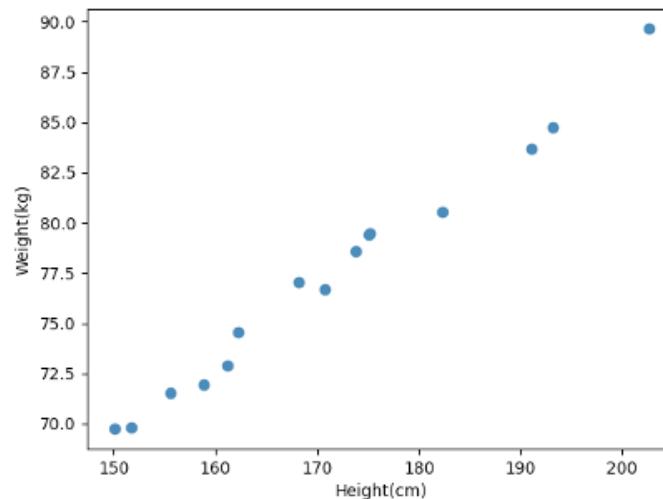
- Regression means learning a **function** that captures the “trend” between input and output.
- The output is a **continuous value**.
- This function is used to predict the target values for new inputs.



# Example of a regression problem

- Can we predict people's weight from their height?

Height(cm)	Weight(kg)
150.00686	69.73347
151.64326	69.83261
155.54032	71.55730
158.80535	71.92875
161.17561	72.92118
:	
175.15167	79.48533
182.32900	80.52182
191.11317	83.67998
193.21947	84.72086
202.68705	89.64049



- Visually, there appears to be a trend.
- A reasonable **model** seems to be the class of linear functions (lines).

# Univariate linear regression

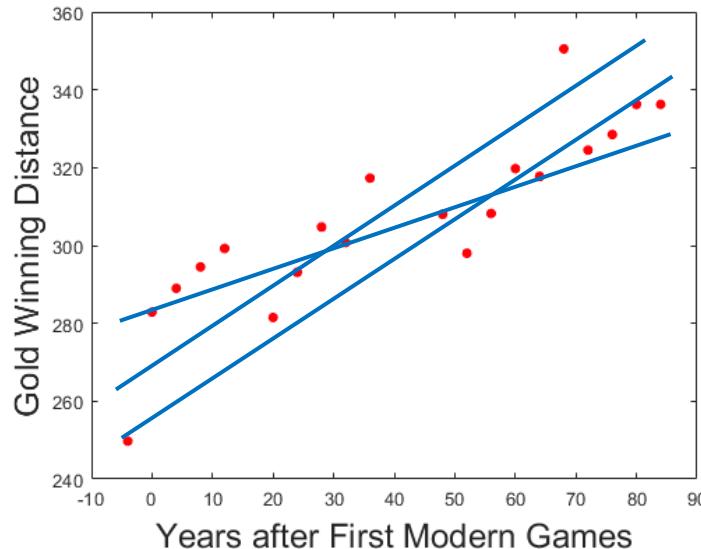
- We have one input attribute (height) – hence the name **univariate**.

$$y = f(x; w_0, w_1) = w_1 x + w_0$$

- 
- Any line is described by this equation by specifying values for  $w_1$  and  $w_0$ .  
dependent variable      free parameters      independent variable



# Our goal: find the “best” line



- Which is the “best” line? That captures the trend in the data.
- Determine the “best” values for  $w_0$  and  $w_1$ .

# Loss/cost functions

- We need a criterion that tells us how good/bad that line is.
- Such criterion is called a loss function.

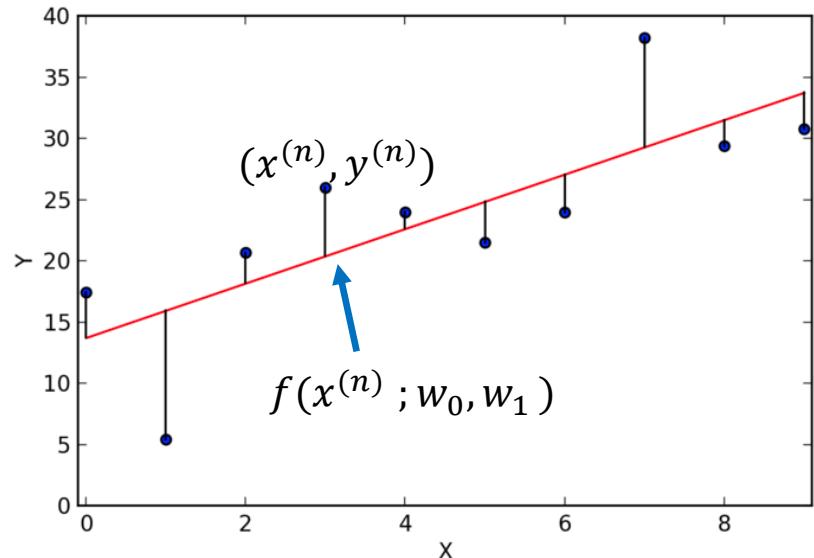
## Terminology

- Loss function = cost function = loss = cost = error function



# We average the losses on all training examples

- For each training example (point)  
 $n = 1, \dots, N$ ,  
The loss on the  $n$ -th point is the  
mismatch/distance between the output of  
the model for this point  
 $f(x^{(n)}; w_0, w_1)$  and the observed target  
 $y^{(n)}$ .
- Average these losses.



# Loss function

- The loss expresses an error, so it must be always non-negative.
- Absolute value loss (L1 loss):

$$L1 = |f(x) - y|$$

- Mean squared error loss (L2 loss):

$$L2 = (f(x) - y)^2$$

$$g(w_0, w_1) = \frac{1}{N} \sum_{n=1}^N (f(x^{(n)}; w_0, w_1) - y^{(n)})^2$$

*Empirical loss  
used by LR*

Loss for the n-th training example

- 0/1 loss:

$$L_{0/1} = 0 \text{ if } f(x) = y, \text{ else } 1$$



# Univariate linear regression

- Given training data

$$(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots (x^{(N)}, y^{(N)})$$

- Fit the model

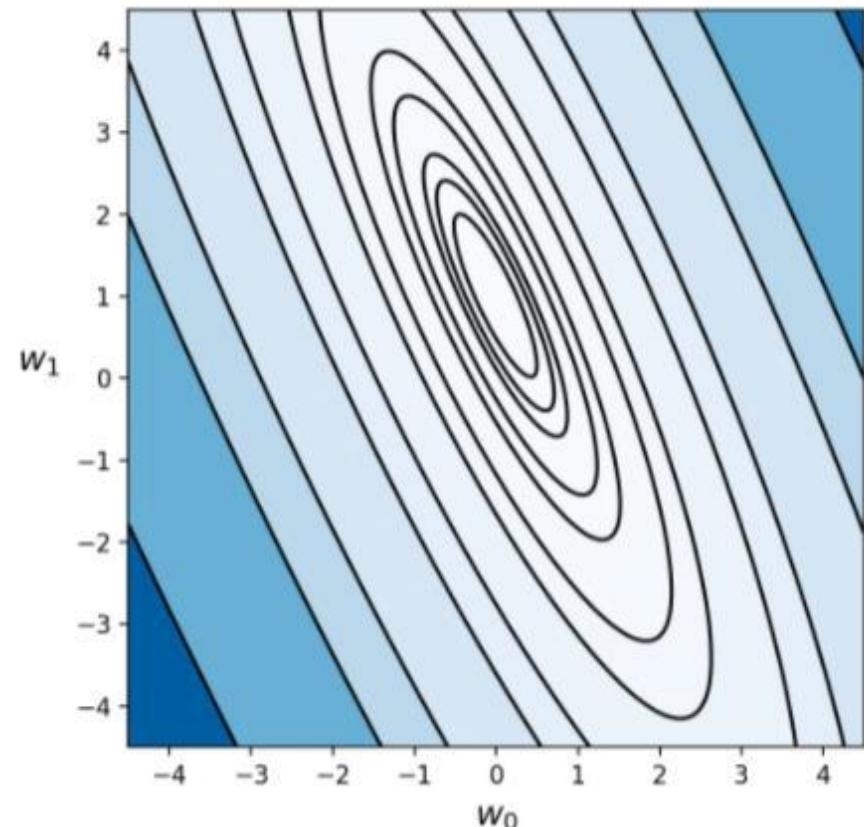
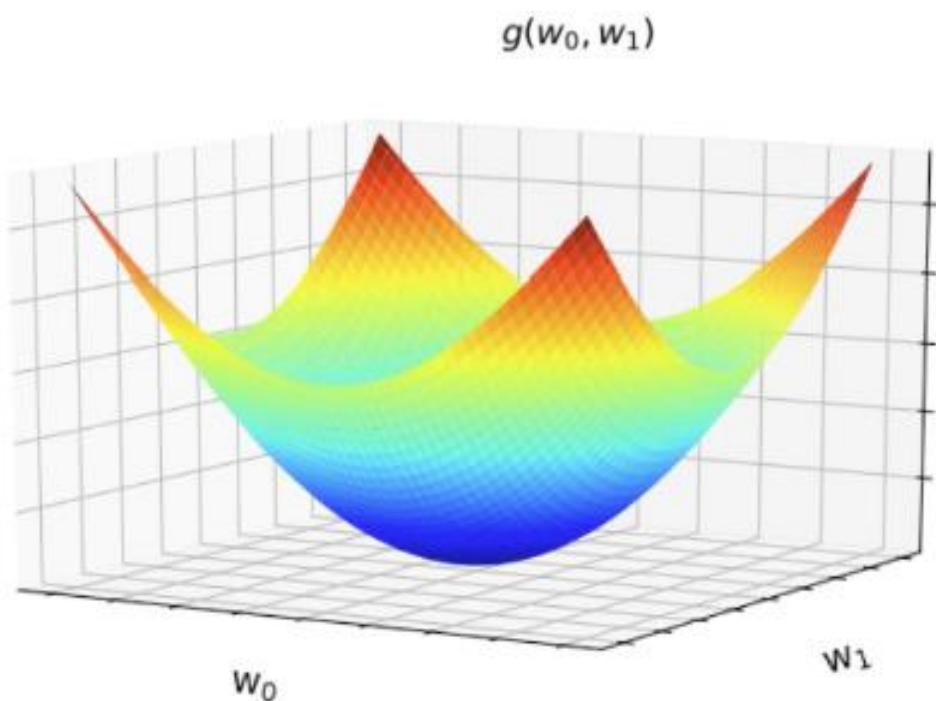
$$y = f(x; w_0, w_1) = w_1 x + w_0$$

- By minimizing the cost function

$$g(w_0, w_1) = \frac{1}{N} \sum_{n=1}^N (f(x^{(n)}; w_0, w_1) - y^{(n)})^2$$

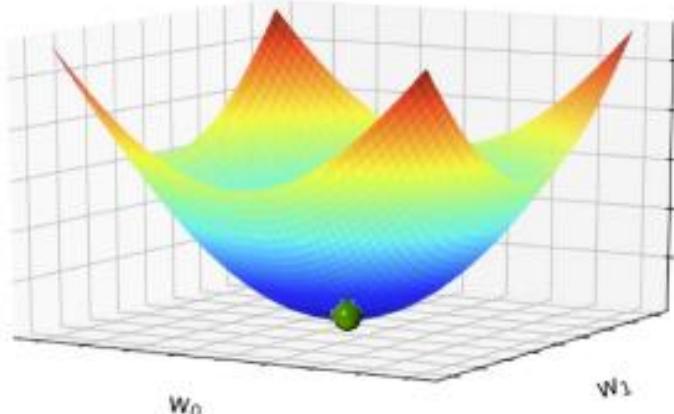
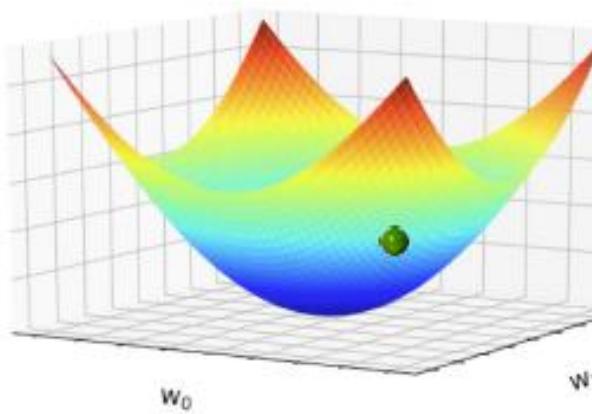
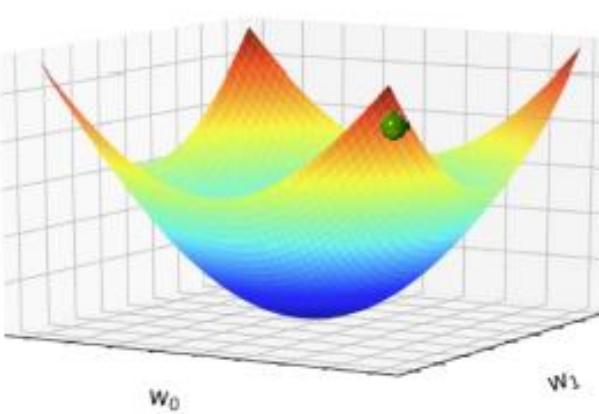
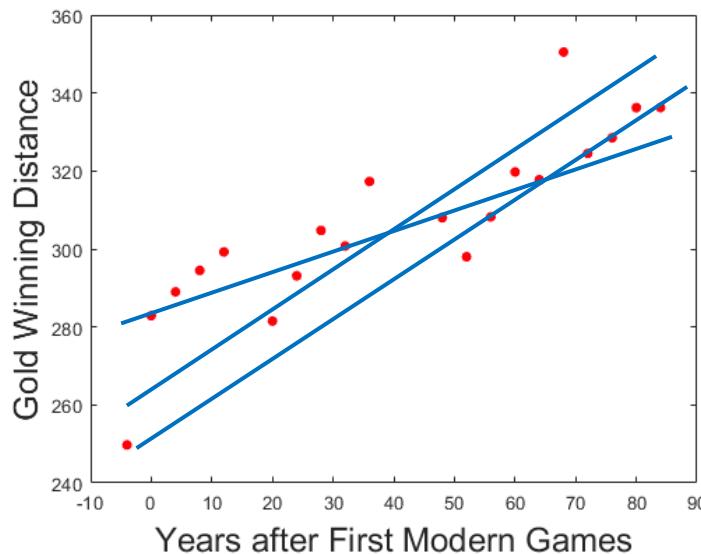


# Cost function depends on the free parameter



# Univariate linear regression

- Every combination of  $(w_0, w_1)$  has an associated cost.
- Key training task: find the ‘best’ values of  $(w_0, w_1)$  such that the cost is minimum.



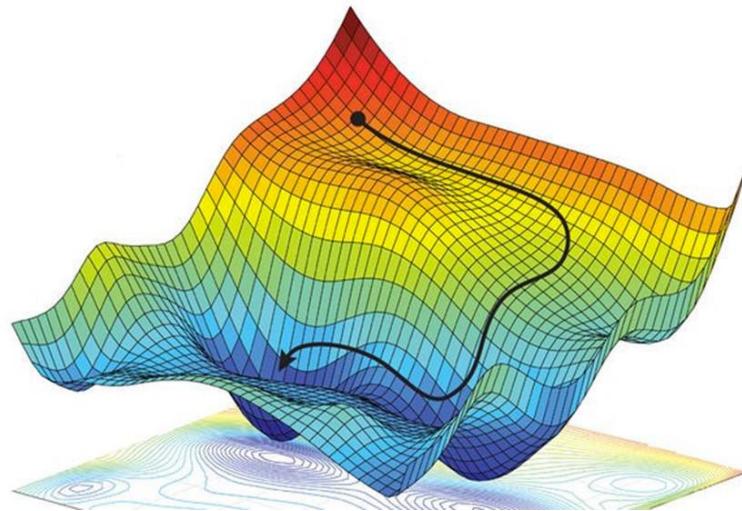
# Gradient Descent



UNIVERSITY OF  
BIRMINGHAM

# Gradient Descent

- A general strategy to minimize cost functions.



# Gradient Descent

- A general strategy to minimize cost functions in ML algorithms.
- Goal: minimize the cost function  $g(w_0, w_1)$

Start at a random point say  $w_0 = 0, w_1 = 0$

Repeat until no change occurs

    Update  $w_0, w_1$  by taking

    a small step in the direction of the steepest descent of cost

Return  $w_0, w_1$ .



# Gradient Descent – More General...

- Goal: minimize the cost function  $g(\mathbf{w})$ , where  $\mathbf{w} = (w_0, w_1, \dots)$

Input:  $\alpha > 0$

Initialise  $\mathbf{w}$ . //at 0 or some random value

Repeat until convergence

$$\mathbf{w} := \mathbf{w} - \alpha \nabla g(\mathbf{w})$$

Return  $\mathbf{w}$ .



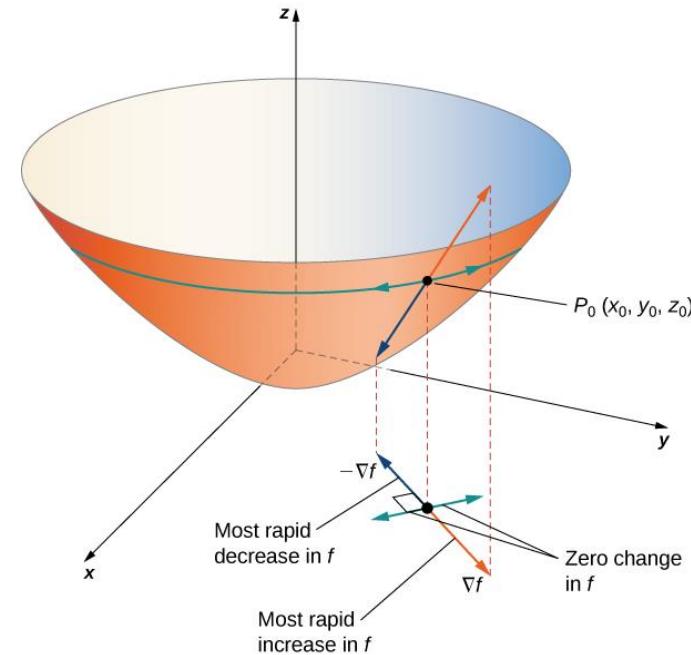
Learning rate  
or step size,  
e.g. 0.01

Gradient or steepest  
direction



UNIVERSITY OF  
BIRMINGHAM

# Back to Univariate Linear Regression



Back to two dimensional function  $g(w_0, w_1)$ :

- The vector of partial derivatives is called the **gradient vector**.

$$\nabla g(\mathbf{w}) = \begin{pmatrix} \frac{\partial g}{\partial w_0} \\ \frac{\partial g}{\partial w_1} \end{pmatrix}, \text{ where } \mathbf{w} = \begin{pmatrix} w_0 \\ w_1 \end{pmatrix}$$

- Recall: Partial derivative with respect to one variable is the ordinary derivative of the function by treating the others as constants.
- The negative of the gradient  $-\nabla g$  evaluated at a location  $(\hat{w}_0, \hat{w}_1)$  gives us the direction of the steepest descent from that location.
- We take a small step in that direction – using the learning rate  $\alpha$ .



# Applying GD to solve univariate linear regression

- Recall: we aim to minimizing the cost function

$$g(w_0, w_1) = \frac{1}{N} \sum_{n=1}^N ((w_1 x^{(n)} + w_0) - y^{(n)})^2$$

- Using the chain rule, we have \*:

$$\frac{\partial g}{\partial w_0} = \frac{2}{N} \sum_{n=1}^N ((w_1 x^{(n)} + w_0) - y^{(n)})$$

$$\frac{\partial g}{\partial w_1} = \frac{2}{N} \sum_{n=1}^N ((w_1 x^{(n)} + w_0) - y^{(n)}) x^{(n)}$$



# Algorithm for univariate linear regression using GD

Input:  $\alpha > 0$ , training set  $\{(x^{(n)}, y^{(n)}): n = 1, 2 \dots N\}$

Initialise  $w_0 = 0, w_1 = 0$

Repeat // (first layer loop)

    for  $n = 1, 2 \dots N$  // second layer loop, more efficient to update after each data point

$$w_0 := w_0 - \alpha((w_1 x^{(n)} + w_0) - y^{(n)})$$

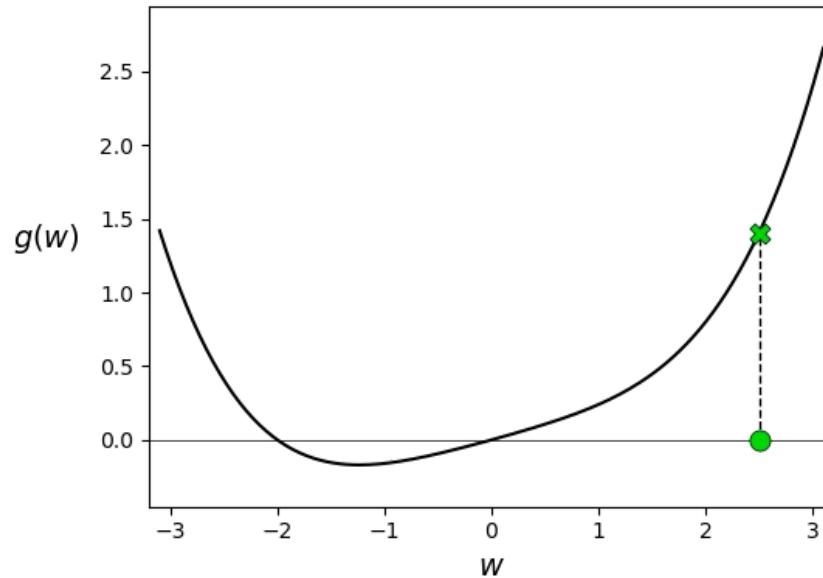
$$w_1 := w_1 - \alpha((w_1 x^{(n)} + w_0) - y^{(n)})x^{(n)}$$

Until change in cost remains below a very small threshold

Return  $w_0, w_1$ .



# Demo Example for Gradient Descent

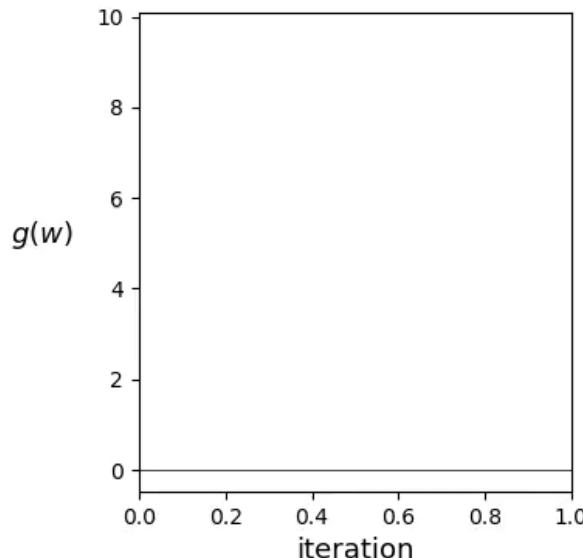
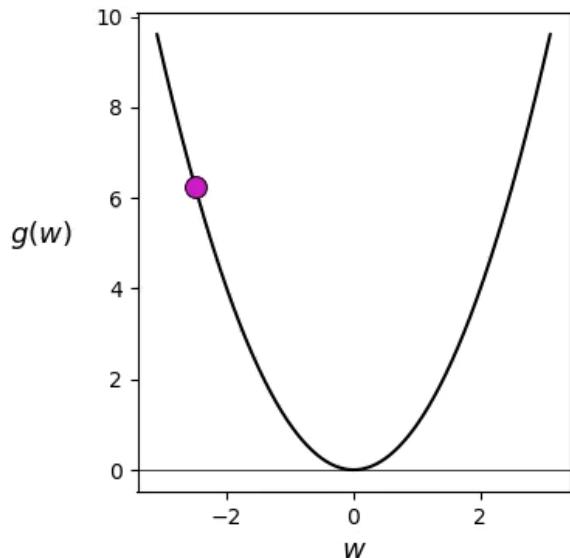


UNIVERSITY OF  
BIRMINGHAM

[https://github.com/jermwatt/machine\\_learning\\_refined/tree/gh-pages/notes/pages/notes/3\\_First\\_order\\_methods/videos/animation\\_6.mp4](https://github.com/jermwatt/machine_learning_refined/tree/gh-pages/notes/pages/notes/3_First_order_methods/videos/animation_6.mp4)

# Effect of the learning rate

Whether or not we descend in the function when taking this step depends completely on how far along it we travel.

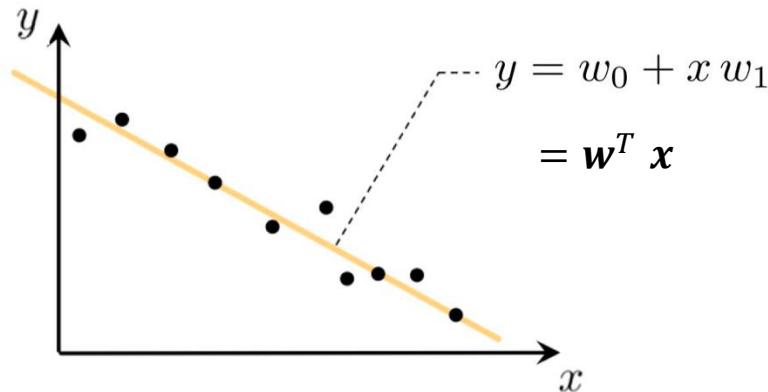


# So far, univariate linear regression

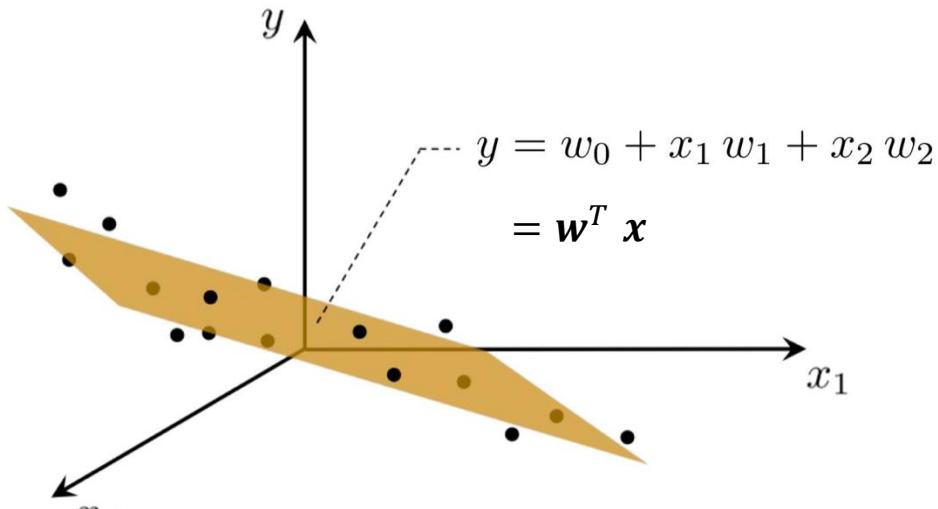
- One input variable
- Assume linear function



# Multivariate linear regression



$$\mathbf{w} = \begin{pmatrix} w_0 \\ w_1 \end{pmatrix}, \mathbf{x} = \begin{pmatrix} 1 \\ x \end{pmatrix}$$



$$\mathbf{w} = \begin{pmatrix} w_0 \\ w_1 \\ w_2 \end{pmatrix}, \mathbf{x} = \begin{pmatrix} 1 \\ x_1 \\ x_2 \end{pmatrix}$$

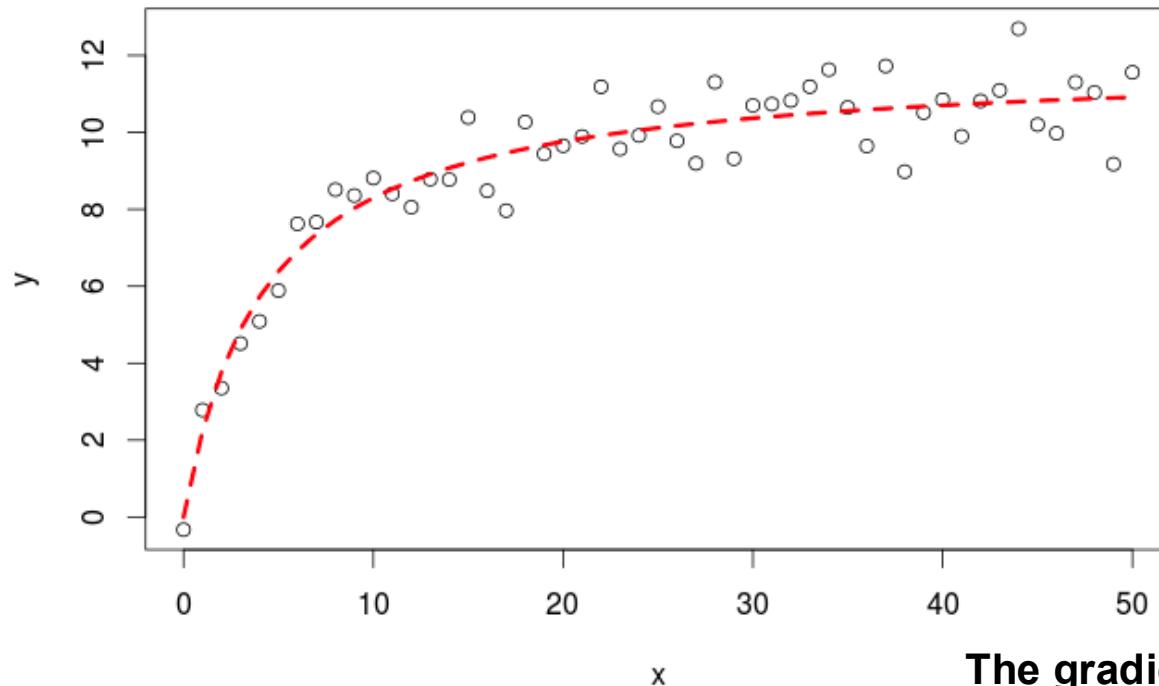
The gradient remains:

$$\nabla g(\mathbf{w}) = 2(\mathbf{w}^T \mathbf{x}^{(n)} - y^{(n)}) \mathbf{x}^{(n)}$$



# Univariate nonlinear regression

$$y = w_0 + w_1 x$$



$$\mathbf{w} = \begin{pmatrix} w_0 \\ w_1 \\ w_2 \\ \dots \\ w_m \end{pmatrix}, \mathbf{x} = \begin{pmatrix} 1 \\ x \\ x^2 \\ \dots \\ x^m \end{pmatrix}$$

This is an  $m$ -th order polynomial regression model.

The gradient remains:

$$\nabla g(\mathbf{w}) = 2(\mathbf{w}^T \mathbf{x}^{(n)} - y^{(n)}) \mathbf{x}^{(n)}$$

# Advantages of vector notation

- Vector notation is more concise.
- With the vectors  $w$  and  $x$  populated appropriately (and differently in each case, as on the previous 2 slides), these models are still linear in the parameter vector.
- The cost function is the L2 as before.
- The gradient remains:

$$\nabla g(w) = 2(w^T x^{(n)} - y^{(n)})x^{(n)}$$

- Ready to be plugged into the general gradient descent algorithm.



# Q/A

**Office Hour and Dropin Sessions**  
**See Canvas module homepage**

**Figures and animations referred to:**

Jeremy Watt et al. Machine Learning Refined. Cambridge University Press, 2020.

[https://github.com/jermwatt/machine\\_learning\\_refined](https://github.com/jermwatt/machine_learning_refined)

